

bpfbox: Simple Precise Process Confinement in eBPF

William Findlay Anil Somayaji David Barrera

Carleton University
`will@ccsl.carleton.ca`

October 18, 2020

Outline of Talk

What is eBPF?

Motivation

Architecture

Policy

Performance

Conclusion

What is eBPF?

eBPF in the Beginning

eBPF \equiv **E**xtended **B**erkley **P**acket **F**ilter...

- ▶ But it has little to do with Berkley, packets, or filtering nowadays

So then **what is eBPF?**

- ▶ A major re-write of the Linux BPF engine
 - ▶ Alexei Starovoitov and Daniel Borkman
- ▶ Merged into the Linux kernel in 2014
- ▶ The point was fine-grained, cross-layer **system introspection**


eBPF in 2020



Motivation

The Status Quo

- ▶ Existing process confinement mechanisms are **complex**



- ▶ Existing process confinement mechanisms are **difficult to use**



- ▶ Can we do any better?

Stakeholders as Policy Authors

- ▶ Security experts define the policy



- ▶ Application authors and **packagers** define the policy



- ▶ End users define the policy

???

eBPF Changes the Game

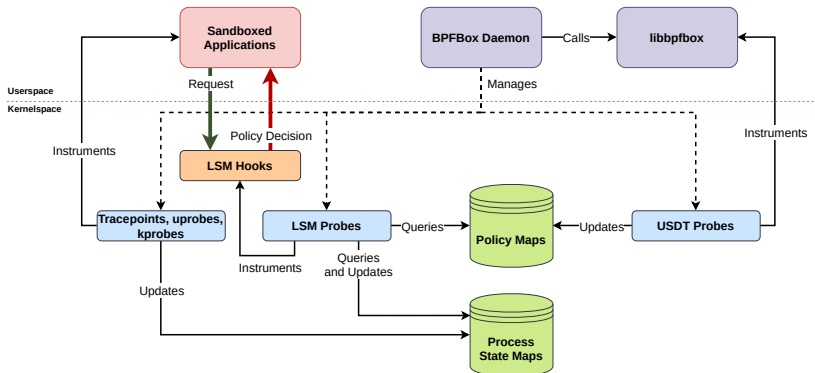
TODO

Architecture

bpfbox Architecture

- ▶ TODO: Python3 bcc
- ▶ TODO: KRSI
- ▶ TODO: Lines of userspace code
- ▶ TODO: Lines of kernelspace code
- ▶ TODO: Compare w/ SELinux, AppArmor

bpffbox Architecture



Policy

Policy at the Function Call Level

```
#![profile /sbin/mylogin]

#[func check_password]
#[func add_user]
#[allow] {
    read("/etc/passwd")
    read("/etc/shadow")
}

#[func add_user]
#[allow] {
    append("/etc/passwd")
    append("/etc/shadow")
}
```

Performance

Performance

TODO

Conclusion

Acknowledgements

TODO

Contributions

- ▶ First full policy enforcement engine written in eBPF
- ▶ Integration of **userspace** and **kernelspace** state with **LSM layer enforcement**