

bpfbox: Simple Precise Process Confinement in eBPF

William Findlay Anil Somayaji David Barrera

Carleton University
`will@ccsl.carleton.ca`

October 18, 2020

Outline of Talk

What is eBPF?

Motivation

bpfbox Architecture

bpfbox Policy

Performance Evaluation

Conclusion

What is eBPF?

eBPF in the Beginning

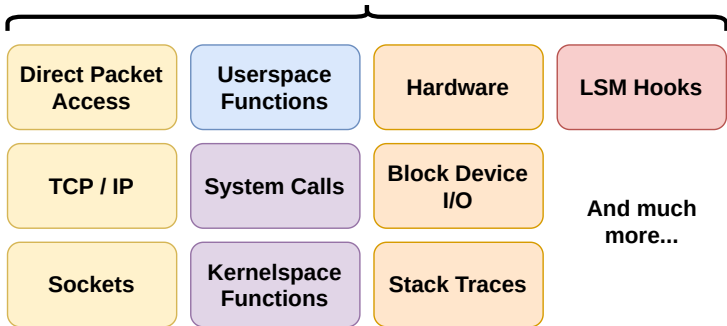
eBPF \equiv **E**xtended **B**erkley **P**acket **F**ilter...

- ▶ But it has little to do with Berkley, packets, or filtering nowadays

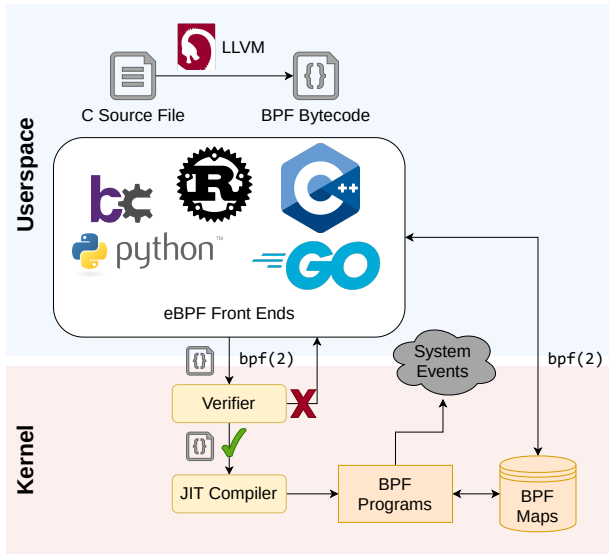
So then **what is eBPF?**

- ▶ A major re-write of the Linux BPF engine
 - ▶ Alexei Starovoitov and Daniel Borkman
- ▶ Merged into the Linux kernel in 2014
- ▶ The point was fine-grained, cross-layer **system introspection**

What Can eBPF Do?



How eBPF Works



eBPF in 2020


eBPF is now **more than just an observability tool.**

- ▶ TODO

Motivation

The Status Quo

- ▶ Existing process confinement mechanisms are **complex**



- ▶ Existing process confinement mechanisms are **difficult to use**



- ▶ Can we do any better?

Stakeholders as Policy Authors

- ▶ Security experts define the policy



- ▶ Application authors and **packagers** define the policy



- ▶ End users define the policy

???

eBPF Changes the Game

TODO

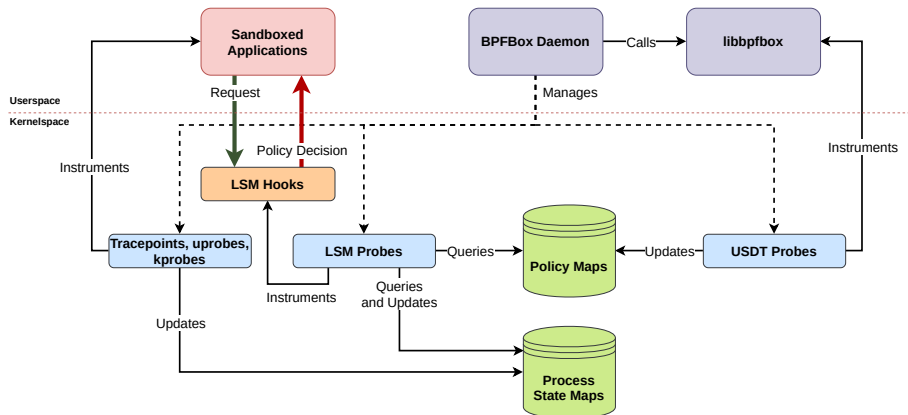
bpfbox

Architecture

bpfbox Architecture

- ▶ TODO: Python3 bcc
- ▶ TODO: KRSI
- ▶ TODO: Lines of userspace code
- ▶ TODO: Lines of kernelspace code
- ▶ TODO: Compare w/ SELinux, AppArmor

bpffbox Architecture



bpfbox Policy

Policy at the Function Call Level

```
#![profile /sbin/mylogin]
```

```
#[func check_password]
```

```
#[func add_user]
```

```
#[allow] {  
    read("/etc/passwd")  
    read("/etc/shadow")  
}
```

```
#[func add_user]
```

```
#[allow] {  
    append("/etc/passwd")  
    append("/etc/shadow")  
}
```


Performance Evaluation

Performance

TODO

Conclusion

Acknowledgements

Special thanks to:

- ▶ **Alexei Starovoitov** and **Daniel Borkman** (creators of eBPF)
- ▶ **K.P. Singh** (creator of KRSI)
- ▶ Fellow **bcc contributors** (an awesome eBPF framework)
- ▶ Anonymous **CCSW'2020 reviewers** (valuable feedback)

This work was supported by NSERC through a Discovery Grant.

Contributions

- ▶ First full policy **enforcement engine** written in eBPF
- ▶ Integration of **userspace** and **kernelspace** state with **LSM layer enforcement**
- ▶ A simple policy language for **ad hoc process confinement**
 - ▶ But with optional complexity for **fine-grained protection**



Check out the project on GitHub!