

# bpfbbox: Simple Precise Process Confinement with eBPF

**William Findlay**

Anil Somayaji

David Barrera

Carleton University

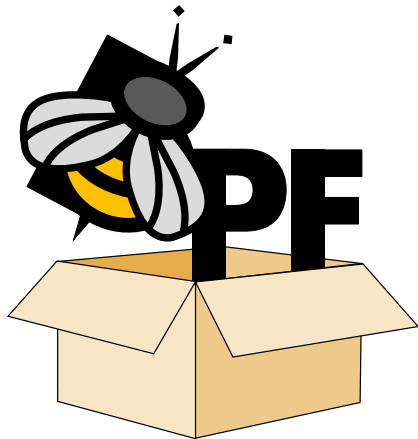
[will@ccsl.carleton.ca](mailto:will@ccsl.carleton.ca)

November 9, 2020



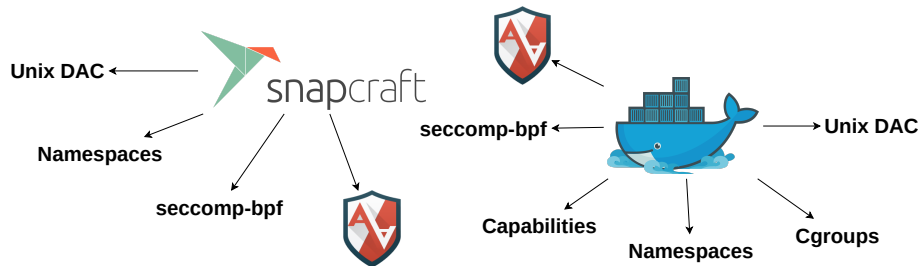
# bpfbbox at a Glance

- ▶ bpfbbox is a novel **process confinement mechanism** for Linux
  - ▶ Using a new Linux technology called eBPF
- ▶ Users write per-application policy in a **simple policy language**
- ▶ Policy is enforced by attaching **eBPF programs** to **LSM hooks**
  - ▶ Integrates cross-layer state into policy decisions



# Motivation

- ▶ Existing process confinement mechanisms are **complex**



- ▶ Existing process confinement mechanisms are **difficult to use**



SELinux



AppArmor



TOMOYO

- ▶ Can we do any better?

# eBPF Changes the Game

eBPF enables:

- ▶ Fine-grained system introspection
- ▶ Integration of **cross-layer state** with policy enforcement
- ▶ Rapid prototyping
- ▶ Safe production deployment of new security solutions

We have an opportunity to **rethink process confinement** from the ground up.

# eBPF in the Beginning

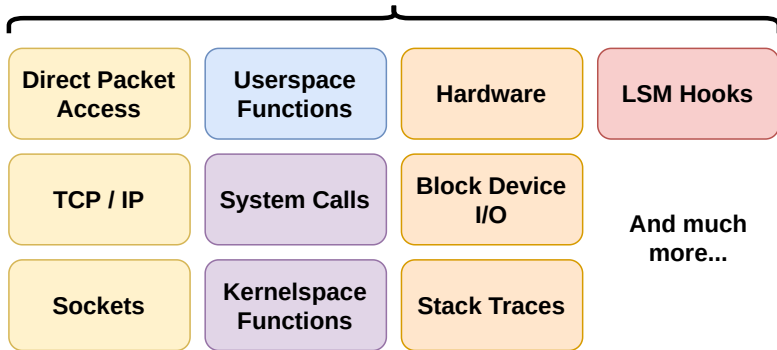
eBPF  $\equiv$  **E**xtended **B**erkley **P**acket **F**ilter...

- ▶ But it has little to do with Berkley, packets, or filtering nowadays
- ▶ The name BPF is preserved for historical reasons

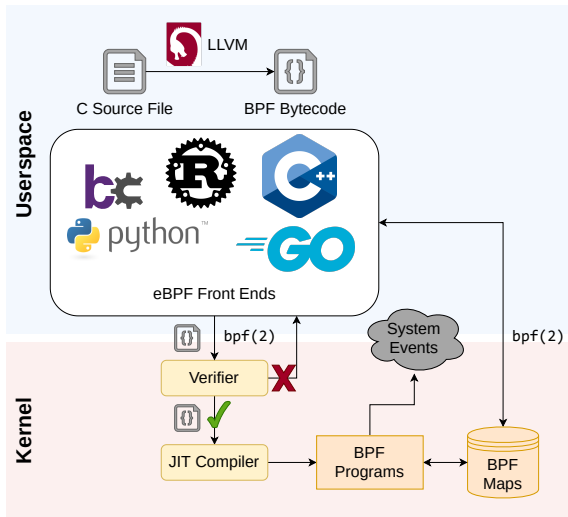
So then **what is eBPF?**

- ▶ A major re-write of the Linux BPF engine
  - ▶ Alexei Starovoitov and Daniel Borkmann
- ▶ Merged into the Linux kernel in 2014
- ▶ The point was fine-grained, cross-layer **system introspection**

# What Can eBPF Do?



# How eBPF Works



# eBPF in 2020

eBPF is now **more than just an observability tool**.

- ▶ eBPF provides a **safe, efficient, and flexible** way for privileged users to extend the kernel
- ▶ eBPF turns Linux into a **programmable kernel**


Linux 5.7 → KRSI (**K**ernel **R**untime **S**ecurity **I**nstrumentation)

- ▶ Attach BPF programs to LSM hooks
- ▶ Make security decisions and generate audit logs with eBPF




# bpffbox Implementation

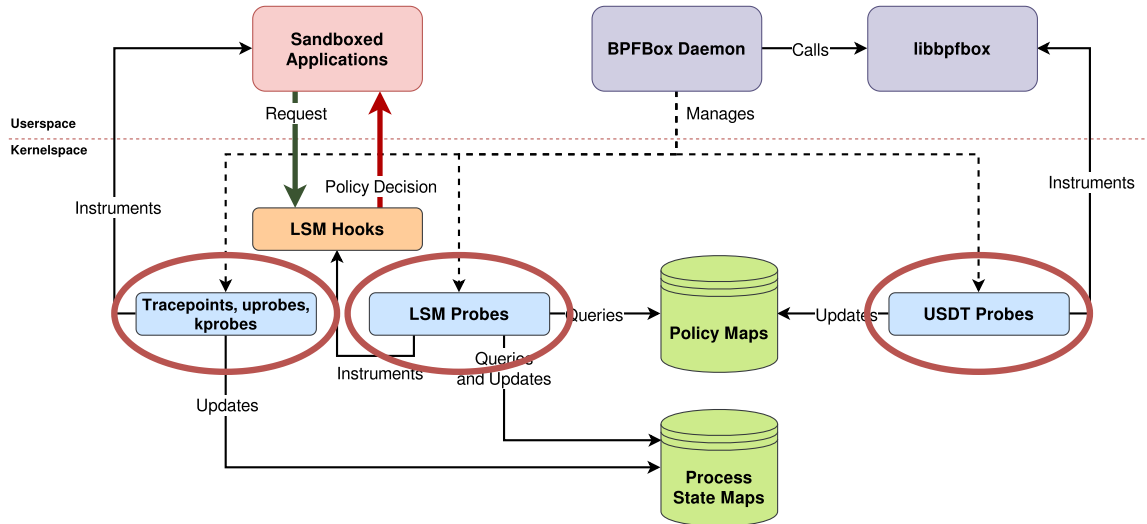
- ▶ Userspace daemon using the Python3 bcc framework
- ▶ Kernelspace components are all eBPF
  - ▶ LSM probes (KRSI), kprobes, uprobes
  - ▶ Under 2000 source lines of kernelspace code
- ▶ Thanks to eBPF, bpffbox is **light-weight**, **flexible**, and **production-safe**
  - ▶ Works out of the box on any vanilla Linux kernel  $\geq 5.8$



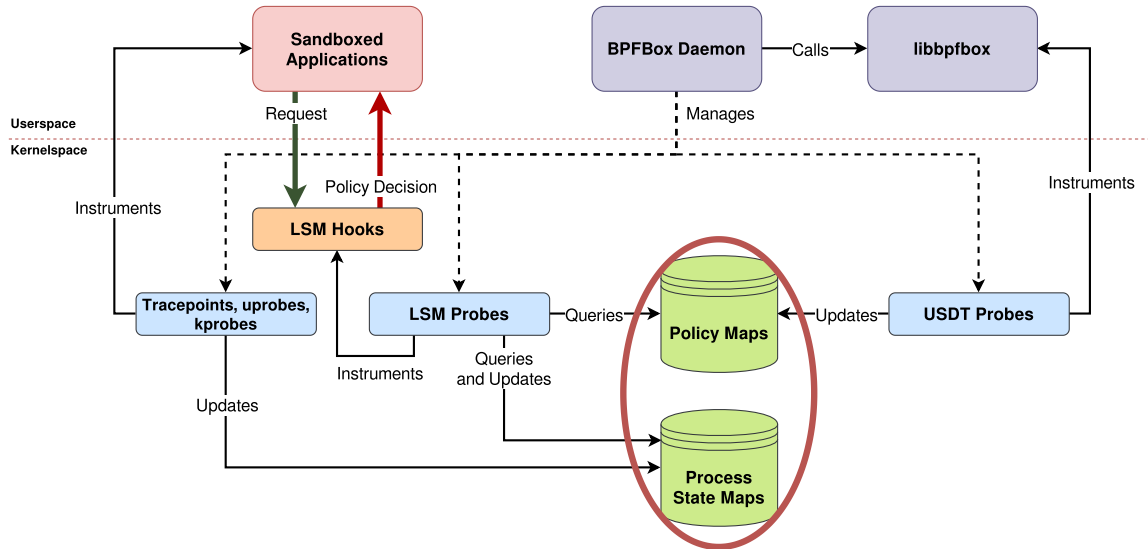
# bpfbox Architecture



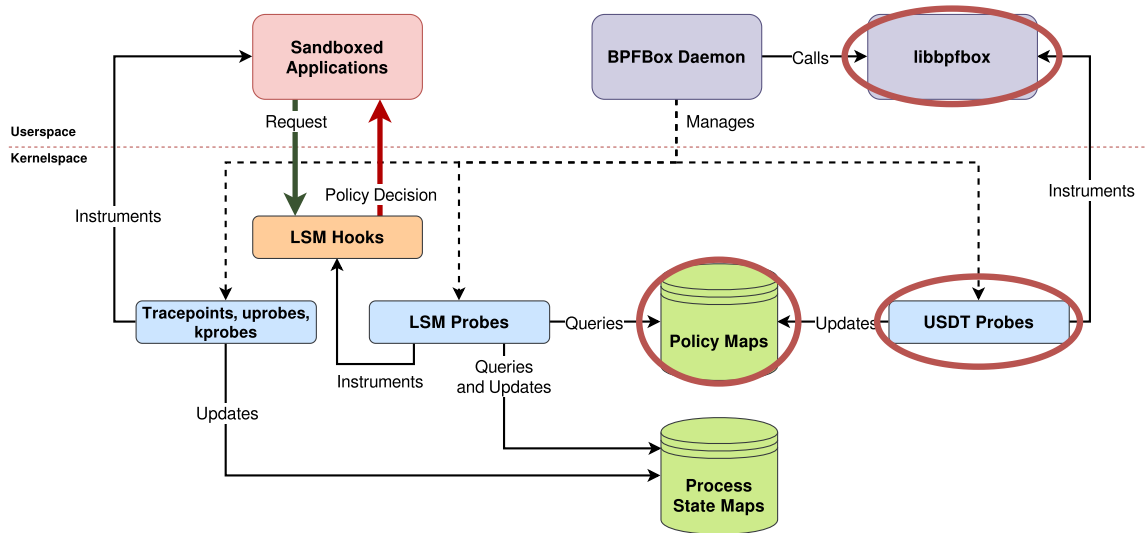
# bpfbox Architecture



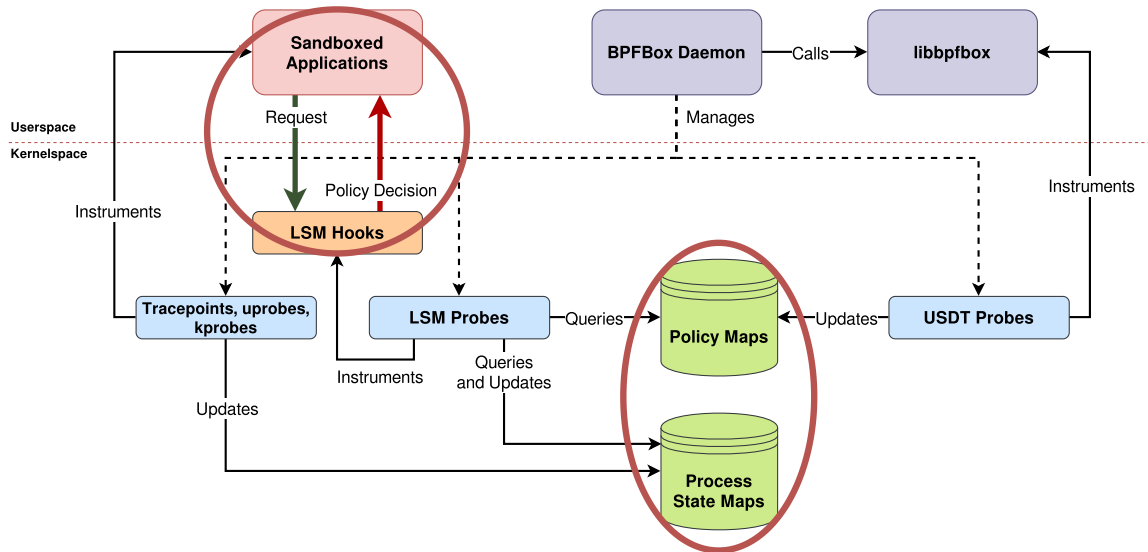
# bpfbox Architecture



# bpfbox Architecture



# bpfbox Architecture



# Policy Design Goals

## 1. Simplicity

- ▶ Policy should be simple enough for ad hoc confinement

## 2. Application transparency

- ▶ Policy should not require changes to the confined application

## 3. Flexibility

- ▶ Policy should offer optional layers of granularity

## 4. Security

- ▶ Policy should follow the principle of least privilege
- ▶ It should be difficult to write an insecure policy

# Rules and Directives

*Rules* specify access to system objects:

- ▶ `fs(file, access)`
- ▶ `net(socket, access)`
- ▶ `signal(prog, sig)`
- ▶ etc.

*Directives* augment blocks of rules:

- ▶ `#[directive]` syntax
- ▶ Specify **actions to be taken** on a block of rules
- ▶ Add **additional context** to a block of rules



# Taints and Transitions

- ▶ `#[taint]` → Start confinement
- ▶ `#[transition]` → Switch profiles on `execve`

```
#![profile "/bin/mywebdaemon"]

#[taint] {
    net(inet, any)
    net(inet6, any)
}

/* ... */

#[transition] {
    fs("/bin/myhelper", getattr|read|exec)
}
```

# Policy at the Function Call Level

- ▶ `#[func "foo"]` → Apply rules only within a call to `foo()`
- ▶ `#[kfunc "foo"]` → Same thing, but for kernel functions

```
#[profile "/sbin/mylogin"]

#[func "check_password"]
#[allow] {
    fs("/etc/passwd", read)
    fs("/etc/shadow", read)
}

#[func "add_user"]
#[allow] {
    fs("/etc/passwd", read|append)
    fs("/etc/shadow", read|append)
}

/* ... */
```

# Performance Evaluation

## Methodology

- ▶ Phoronix Test Suite OSBench
  - ▶ Measures basic OS functionality
  - ▶ (spawning processes, memory allocations, etc.)
- ▶ Phoronix Test Suite Apache
  - ▶ Benchmark Apache `httpd` packets per second
- ▶ Kernel compilation benchmarks
  - ▶ Measure Linux kernel compilation performance
  - ▶ Heavy workload, spawning lots of processes

# Performance Evaluation

## Methodology

Two modes of operation for each test.

- ▶ Passive mode
  - ▶ bpfbox and AppArmor instrument hooks, but do not enforce or audit
  - ▶ Test lowest possible overhead
- ▶ Complaining mode
  - ▶ bpfbox and AppArmor complain about (log) every security-sensitive operation
  - ▶ Test worst case overhead

# Performance Evaluation

## Results

- ▶ Phoronix OSBench
  - ▶ Passive: bpfbox is **roughly equivalent** to AppArmor
  - ▶ Complaining: bpfbox performs **significantly better** than AppArmor
- ▶ Phoronix Apache
  - ▶ bpfbox and AppArmor are **roughly equivalent**
- ▶ Kernel compilation
  - ▶ Passive: bpfbox is **roughly equivalent** to AppArmor
  - ▶ Complaining: bpfbox performs **better in kernelspace** overhead and **worse in userspace** overhead

# Contributions

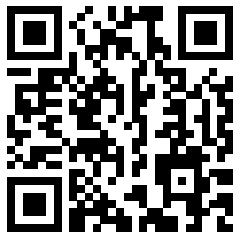
- ▶ First **policy enforcement engine** written in **eBPF**
- ▶ Integration of **userspace** and **kernelspace** state with **LSM layer enforcement**
- ▶ A simple policy language for **ad hoc process confinement**
  - ▶ But with optional complexity for **fine-grained protection**

# Acknowledgements

Special thanks to:

- ▶ **Alexei Starovoitov** and **Daniel Borkmann** (creators of eBPF)
- ▶ **K.P. Singh** (creator of KRSI)
- ▶ Fellow **bcc contributors** (an awesome eBPF framework)
- ▶ Anonymous **CCSW'2020 reviewers** (valuable feedback)

This work was supported by NSERC through a Discovery Grant.



[github.com/willfindlay/bpfbox](https://github.com/willfindlay/bpfbox)

Check out the project on GitHub!